

# A Causal Inference Perspective on Therapist Effects—Online Supplement

Kristoffer Magnusson

2023-06-16

## Table of contents

Introduction . . . . .	1
Therapist-Outcome Confounding Simulation . . . . .	2
A More Intuitive Interpretation of Therapist Effects . . . . .	12
Confidence Intervals . . . . .	24

## Introduction

This online supplement for *A Causal Inference Perspective on Therapist Effects* (Magnusson, 2023, in preparation) contains the code for all calculations, simulations, and figures presented in the article.

- [Preprint](#)
- [GitHub repository](#)
- [OSF repository](#)
- [Interactive visualization](#)

An HTML version of this document can also be viewed at <https://rpsychologist.github.io/causal-therapist-effects-paper>

```
knitr::opts_chunk$set(  
  message = FALSE,  
  warning = FALSE,  
  cache = TRUE  
)
```

```

library(ggplot2)
library(readr)
library(dplyr)
library(tidyr)
library(purrr)
library(powerlmm)
library(knitr)
library(parallel)
library(lme4)
library(lmerTest)
library(svglight)
library(brms)
library(parallel)

```

## Therapist-Outcome Confounding Simulation

The functions below were used to simulate a data set with therapist effects and a therapist-outcome confounder.

```

#' Get the mean difference between prognostic groups
#'
#' @param sd the SD added to the therapist level due to the confounder
#' @param n2 the number of patients per therapist
#'
#' @returns the mean difference between the two binary groups
#' @examples
#' n2 <- 50
#' M <- solve_for_mean(0.075, n2)
#' # we solve for M in this
#' sd(c(rep(0, n2), rep(M, n2)))
solve_for_mean <- function(sd, n2) {
  2 * sqrt(sd^2 * (n2 * 2 - 1) / (n2 * 2))
}

#' Get the therapist SD from the ICC
#'
#' @param icc the ICC
#' @param sd_error the error SD
#'
#' @returns the random therapist SD

```

```

get_therapist_sd_from_icc <- function(icc, sd_error) {
  (sqrt(icc) * sd_error) / (sqrt(1 - icc))
}

#' Simulate confounded therapist effects
#'
#' @param n1
#' @param n2
#' @param sd_therapist
#' @param sd_therapist_confounding
#' @param sd_error
#' @param ATE
#'
#' @returns a data.frame
simulate_therapist_effect_confounding <- function(
  n1,
  n2,
  sd_therapist,
  sd_therapist_confounding,
  sd_error,
  ATE,
  ...) {
  tot_n <- 2 * n1 * n2
  b_pre_prognosis <- solve_for_mean(
    sd = sd_therapist_confounding,
    n2 = n2
  )
  Z <- rep(c(0, 1), each = n1 * n2)
  therapist_cc <- 1:n2
  therapist_tx <- (n2 + 1):(n2 * 2)
  therapist_effect_cc <- rnorm(n2, 0, sd_therapist)
  therapist_effect_tx <- rnorm(n2, 0, sd_therapist)
  therapist_effects <- c(therapist_effect_cc, therapist_effect_tx)
  error <- rnorm(tot_n, 0, sd_error)
  d <- data.frame(
    Z,
    pre_prognosis = rbinom(tot_n, 1, 0.5),
    therapist = NA,
    therapist_random = NA
  )
  d$therapist <- ifelse(

```

```

Z == 0,
# control
ifelse(
  d$pre_prognosis == 0,
  sample(
    x = 1:(n2 / 2),
    size = n1 * n2 / 2,
    replace = TRUE
  ),
  sample(
    x = (n2 / 2 + 1):(n2),
    size = n1 * n2 / 2,
    replace = TRUE
  )
),
# treatment
ifelse(
  d$pre_prognosis == 0,
  sample(
    x = 1:(n2 / 2),
    size = n1 * n2 / 2,
    replace = TRUE
  ),
  sample(
    x = (n2 / 2 + 1):(n2),
    size = n1 * n2 / 2,
    replace = TRUE
  )
) + n2
)
d$therapist_random[d$Z == 0] <- sample(
  x = therapist_cc,
  size = n1 * n2,
  replace = TRUE
)
d$therapist_random[d$Z == 1] <- sample(
  x = therapist_tx,
  size = n1 * n2,
  replace = TRUE
)
# non-random allocation of therapists

```

```

d$therapist_effect <- therapist_effects[d$therapist]
d$y <- 10 +
  d$Z * ATE +
  d$therapist_effect +
  error +
  d$pre_prognosis * b_pre_prognosis
# random allocation of therapists
d$therapist_effect_random <- therapist_effects[d$therapist_random]
d$y_random <- 10 +
  d$Z * ATE +
  d$therapist_effect_random +
  error +
  d$pre_prognosis * b_pre_prognosis
# required by powerlmm
d$time <- 0

d
}

```

To run the simulation we pass the simulation function `simulate_therapist_effect_confounding` and use a custom `powerlmm` model to setup all the parameters.

```

ds <- study_design(custom = TRUE)
ICC <- 0.05
sd_error <- 1.5
sd_therapist <- sqrt((ICC * sd_error^2) / (1 - ICC))
sd_therapist^2 / (sd_therapist^2 + sd_error^2)

```

[1] 0.05

```

p <- powerlmm:::study_parameters.plcp_design_custom(
  design = ds,
  n1 = 20,
  n2 = 10,
  sd_therapist = sd_therapist,
  sd_error = sd_error,
  ATE = 0,
  data_gen = simulate_therapist_effect_confounding
)
p$sd_therapist_confounding <- get_therapist_sd_from_icc(

```

```

    icc = 0.075,
    sd_error = p$sd_error
  )
  tot_sd <- sqrt(
    p$sd_error^2 + p$sd_therapist^2 + p$sd_therapist_confounding^2
  )
  p$ATE <- 0.5 * tot_sd
  p$tot_n <- 2 * p$n1 * p$n2
  p$df <- 2 * p$n2 - 2

```

True parameter values

```

p$thetas_FE <- list(
  "(Intercept)" = 10,
  "Z" = p$ATE
)
p$thetas_RE <- list(
  "therapist_(Intercept)" = p$sd_therapist^2,
  "therapist_random_(Intercept)" = p$sd_therapist^2,
  "error" = p$sd_error^2
)

```

We then specify the statistical models

```

f0 <- sim_formula(
  "y_random ~ Z + (1 | therapist_random)",
  test = "Z"
)
f1 <- sim_formula(
  "y_random ~ Z + pre_prognosis + (1 | therapist_random)",
  test = "Z"
)
f2 <- sim_formula(
  "y ~ Z + (1 | therapist)",
  test = "Z"
)
f3 <- sim_formula(
  "y ~ Z + pre_prognosis + (1 | therapist)",
  test = "Z"
)
f4 <- sim_formula(

```

```

      "y ~ Z",
      test = "Z"
    )
    f5 <- sim_formula(
      "y ~ Z + pre_prognosis",
      test = "Z"
    )
    f <- sim_formula_compare(
      "rand" = f0,
      "rand_adj" = f1,
      "confounding" = f2,
      "adjusted" = f3,
      "ignored" = f4,
      "ignored_adjusted" = f5
    )

```

Finally, we run the simulation and save the results.

```

# Load cache if it exists
file_path <- "tmp/simulation.rds"
if (file.exists(file_path)) {
  res <- read_rds(file = file_path)
} else {
  MAX_CORES <- as.numeric(Sys.getenv("MAX_CORES"))
  N_SIM <- as.numeric(Sys.getenv("N_SIM"))
  if (is.na(MAX_CORES)) MAX_CORES <- parallel::detectCores(logical = FALSE) - 1
  cl <- makeCluster(MAX_CORES)
  clusterExport(
    cl,
    c(
      "solve_for_mean",
      "get_therapist_sd_from_icc"
    )
  )
  res <- simulate(
    p,
    nsim = ifelse(is.na(N_SIM), 10000, N_SIM),
    cores = MAX_CORES,
    formula = f,
    cl = cl,
    satterthwaite = TRUE
  )
}

```

```

    )
    stopCluster(cl)
    write_rds(
      res,
      file = file_path,
      compress = "gz"
    )
  }

```

Number of simulations:  $10^4$

## Simulations results

We first summarize the treatment effects, and the result is shown in [Table 1](#).

```

summary(
  res,
  verbose = FALSE,
  para = "Z"
)$summary$summary$FE %>%
  mutate(
    SD_rel_bias = (M_se - SD_est) / SD_est
  ) %>%
  relocate(model, .before = 1) %>%
  select(
    -parameter,
    -Power,
    -Power_bw,
    "Model" = model,
    "Estimate" = M_est,
    "Rel. bias (SE)" = SD_rel_bias,
    "SD(Est.)" = SD_est,
    "Power" = Power_satt,
    "True" = theta
  ) %>%
  kable(digits = 2)

```



Table 1: Simulation results (treatment effects)

	Model	Estimate	True	M_se	SD(Est.)	Power	Rel. bias (SE)
2	rand	0.8	0.8	0.22	0.22	0.93	0.00
21	rand_adj	0.8	0.8	0.21	0.21	0.94	0.00
22	confounding	0.8	0.8	0.29	0.21	0.80	0.35
23	adjusted	0.8	0.8	0.21	0.21	0.94	0.00
24	ignored	0.8	0.8	0.16	0.22	0.99	-0.28
25	ignored_adjusted	0.8	0.8	0.15	0.22	0.99	-0.29

A summary of the estimated ICCs is shown in Table 2.

```
lapply(
  seq_along(res$res),
  function(i) {
    res$res[[i]]$RE %>%
      group_by(sim) %>%
      summarize(ICC = vcov[1] / sum(vcov)) %>%
      ungroup() %>%
      summarize(
        est_mean = mean(ICC),
        est_sd = sd(ICC),
        est_lwr = quantile(ICC, 0.025),
        est_upr = quantile(ICC, 0.975)
      ) %>%
      ungroup() %>%
      mutate(
        model = names(res$res)[[i]],
        parameter = "ICC",
        .before = 1
      )
  }
) %>%
  bind_rows() %>%
  filter(grepl("ignored", model) == FALSE) %>%
  mutate(
    theta = p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2),
    rel_bias = (est_mean - theta) / theta
  ) %>%
  select(
    "Model" = model,
```

```

    "Parameter" = parameter,
    "Estimate" = est_mean,
    "Rel. bias" = rel_bias,
    "SD(Est.)" = est_sd,
    "Est. (2.5%)" = est_lwr,
    "Est. (97.5%)" = est_upr,
    "True" = theta
  ) %>%
  kable(digits = 2)

```

Table 2: Simulation results (ICCs)

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand	ICC	0.05	-0.07	0.03	0.00	0.11	0.05
rand_adj	ICC	0.05	0.00	0.03	0.00	0.12	0.05
confounding	ICC	0.12	1.39	0.04	0.04	0.21	0.05
adjusted	ICC	0.05	-0.01	0.03	0.00	0.12	0.05

Lastly, we summarize the variance components.

```

lapply(
  seq_along(res$res),
  function(i) {
    res$res[[i]]$RE %>%
      group_by(parameter) %>%
      summarize(
        est_mean = mean(vcov),
        est_sd = sd(vcov),
        est_lwr = quantile(vcov, 0.025),
        est_upr = quantile(vcov, 0.975)
      ) %>%
      ungroup() %>%
      mutate(
        model = names(res$res)[[i]],
        .before = 1
      )
  }
) %>%
  bind_rows() %>%

```

```

left_join(
  data.frame(
    parameter = names(p$thetas_RE),
    theta = unlist(p$thetas_RE)
  )
) %>%
mutate(
  parameter = replace(
    parameter,
    grep("therapist", parameter),
    "therapist"
  ),
  theta = case_when(
    parameter == "therapist" ~ p$sd_therapist^2,
    parameter == "error" & model == "rand" ~ p$sd_error^2 +
      p$sd_therapist_confounding^2,
    model == "ignored" ~ p$sd_error^2 + p$sd_therapist^2 +
      p$sd_therapist_confounding^2,
    model == "ignored_adjusted" ~ p$sd_error^2 + p$sd_therapist^2,
    parameter == "error" ~ p$sd_error^2
  ),
  rel_bias = (est_mean - theta) / theta
) %>%
select(
  "Model" = model,
  "Parameter" = parameter,
  "Estimate" = est_mean,
  "Rel. bias" = rel_bias,
  "SD(Est.)" = est_sd,
  "Est. (2.5%)" = est_lwr,
  "Est. (97.5%)" = est_upr,
  "True" = theta
) %>%
kable(digits = 2)

```

Table 3: Simulation results (variance components)

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand	error	2.42	0.00	0.18	2.08	2.78	2.43
rand	therapist	0.12	0.01	0.08	0.00	0.30	0.12

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand_adj	error	2.25	0.00	0.16	1.93	2.58	2.25
rand_adj	therapist	0.12	0.01	0.08	0.00	0.30	0.12
confounding	error	2.25	0.00	0.17	1.93	2.58	2.25
confounding	therapist	0.31	1.61	0.13	0.10	0.59	0.12
adjusted	error	2.25	0.00	0.17	1.93	2.58	2.25
adjusted	therapist	0.12	0.00	0.08	0.00	0.30	0.12
ignored	error	2.53	-0.01	0.19	2.16	2.93	2.55
ignored_adj	error	2.35	-0.01	0.17	2.03	2.70	2.37

## A More Intuitive Interpretation of Therapist Effects

This section presents R code for calculating the overlap measures presented in the manuscript.

### Calculate the overlap between therapist distributions

Overlap can be calculated using the examples presented below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```
# Integration
int_f <- function(x, mu1, mu2, sd1, sd2) {
  f1 <- dnorm(x, mean = mu1, sd = sd1)
  f2 <- dnorm(x, mean = mu2, sd = sd2)
  pmin(f1, f2)
}

cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
# standardize using therapist SD
z <- ATE / p$sd_therapist
2 * pnorm(-abs(z) / 2)
```

```
[1] 0.6547208
```

```
# unstandardized
2 * pnorm(
  -abs(ATE) / 2,
```

```
sd = p$sd_therapist
)
```

[1] 0.6547208

```
# Integrate unstandardized
integrate(
  int_f,
  -Inf,
  Inf,
  mu1 = 0,
  mu2 = ATE,
  sd1 = p$sd_therapist,
  sd2 = p$sd_therapist
)
```

0.6547209 with absolute error < 1.6e-05

```
# cohen's d
2 * pnorm(
  -abs(cohensd) / 2,
  sd = p$sd_therapist / tot_sd
)
```

[1] 0.6547208

```
ICC <- p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2)
2 * pnorm(
  -abs(cohensd) / 2,
  sd = sqrt(ICC)
)
```

[1] 0.6547208

## Plot Overlap

Figure 1 visualizes the overlapping therapist effect distributions.

```

SD <- sqrt(0.05)
mean1 <- 0.2
# create x axis
x_min <- 0 - 3 * SD
x_max <- mean1 + 3 * SD
x <- seq(x_min, x_max, length.out = 2e4)
df_control <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, 0, SD)),
  data.frame("x" = x_max, "y" = 0)
)
df_tx <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, mean1, SD)),
  data.frame("x" = x_max, "y" = 0)
)
poly_overlap <- data.frame(
  "x" = df_control$x,
  "y" = pmin(df_control$y, df_tx$y)
)
# colors
overlap_fill <- "#2980b9"
u3_fill <- "#3498db"
control_fill <- "#7f8c8d"
treatment_fill <- "#2c3e50"
p0 <- ggplot(
  df_tx,
  aes(
    x,
    y,
    fill = "treatment",
  )
) +
  # fill treatment group
  geom_polygon(
    linewidth = 1,
  ) +
  # fill control group
  geom_polygon(
    data = df_control,
    aes(

```

```

        fill = "control"
    ),
    linewidth = 1
) +
# overlap
geom_polygon(
    data = poly_overlap,
    color = NA,
    fill = overlap_fill,
) +
# line treatment
geom_polygon(
    linewidth = 1,
    color = "white",
    alpha = 0.5,
    fill = NA
) +
# line control
geom_polygon(
    data = df_control,
    linewidth = 1,
    color = "white",
    alpha = 0.5,
    fill = NA
) +
geom_vline(
    xintercept = 0,
    linetype = "dotted"
) +
geom_vline(
    xintercept = mean1,
    linetype = "dotted"
) +
annotate(
    geom = "text",
    label = "Control",
    x = 0,
    y = dnorm(0, 0, SD) * 1.1
) +
annotate(
    geom = "text",

```

```

    label = "Treatment",
    x = mean1,
    y = dnorm(mean1, mean1, SD) * 1.1
  ) +
  scale_color_manual(
    values = c(
      "control" = control_fill,
      "treatment" = treatment_fill
    )
  ) +
  scale_fill_manual(
    values = c(
      "control" = control_fill,
      "treatment" = treatment_fill
    )
  ) +
  labs(x = "Therapist effects", y = NULL) +
  theme_minimal() +
  theme(
    legend.position = "none",
    panel.grid.minor.y = element_blank(),
    panel.grid.major.y = element_blank(),
    axis.text.y = element_blank()
  )
p0
ggsave("figures/fig_overlap.svg", width = 8, height = 3)

```






Figure 1: The proportion of therapists from each treatment group that have overlapping effects.

### Cohen's $U_3$

Cohen's  $U_3$  can be calculated using the examples below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```
# standardize using therapist SD
cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
# standardize using therapist SD
z <- ATE / p$sd_therapist
pnorm(z)
```

```
[1] 0.8144533
```

```
# unstandardized effect
pnorm(ATE, sd = p$sd_therapist)
```

```
[1] 0.8144533
```

```
# cohen's d parameterization
pnorm(
  cohensd,
  sd = p$sd_therapist / tot_sd
)
```

```
[1] 0.8144533
```

```
ICC <- p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2)
pnorm(
  cohensd,
  sd = sqrt(ICC)
)
```

```
[1] 0.8144533
```

### Plot Cohen's $U_3$

Figure 2 visualizes Cohen's  $U_3$ .

```
SD <- sqrt(0.05)
ES <- 0.2
mean1 <- ES
# create x axis
x_min <- 0 - 3 * SD
x_max <- mean1 + 3 * SD
x <- seq(x_min, x_max, length.out = 2e4)
df_control <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, 0, SD)),
  data.frame("x" = x_max, "y" = 0)
)
df_tx <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, mean1, SD)),
  data.frame("x" = x_max, "y" = 0)
)
poly_u3 <- rbind(
  data.frame("x" = x_min, "y" = 0),
```

```

    poly_overlap[poly_overlap$x <= 0, ],
    data.frame("x" = 0, "y" = 0)
  )
  # colors
  overlap_fill <- "#2980b9"
  u3_fill <- "#3498db"
  control_fill <- "#7f8c8d"
  treatment_fill <- "#2c3e50"
  # plot
  p0 <- ggplot(
    df_control,
    aes(
      x,
      y,
      fill = "control"
    )
  ) +
    # fill control group
    geom_polygon(
      linewidth = 1,
    ) +
    # fill treatment group
    geom_polygon(
      data = df_tx,
      aes(
        fill = "treatment"
      ),
      linewidth = 1
    ) +
    # overlap
    geom_polygon(
      data = poly_u3,
      color = NA,
      fill = u3_fill
    ) +
    # line control
    geom_polygon(
      linewidth = 1,
      color = "white",
      alpha = 0.5,
      fill = NA
    )

```

```

) +
# line treatment
geom_polygon(
  data = df_tx,
  linewidth = 1,
  color = "white",
  alpha = 0.5,
  fill = NA
) +
geom_vline(
  xintercept = 0,
  linetype = "dotted"
) +
geom_vline(
  xintercept = mean1,
  linetype = "dotted"
) +
annotate(
  geom = "text",
  label = "Control",
  x = 0,
  y = dnorm(0, 0, SD) * 1.1
) +
annotate(
  geom = "text",
  label = "Treatment",
  x = mean1,
  y = dnorm(mean1, mean1, SD) * 1.1
) +
scale_color_manual(
  values = c(
    "control" = control_fill,
    "treatment" = treatment_fill
  )
) +
scale_fill_manual(
  values = c(
    "control" = control_fill,
    "treatment" = treatment_fill
  )
) +

```

```

labs(x = "Therapist effects", y = NULL) +
theme_minimal() +
theme(
  legend.position = "none",
  panel.grid.minor.y = element_blank(),
  panel.grid.major.y = element_blank(),
  axis.text.y = element_blank()
)
p0
ggsave("figures/fig_u3.svg", width = 8, height = 3)

```




Figure 2: A visualization of Cohen's U3, the proportion of the therapists in the treatment group with causal effects above the average therapist in the control group

### Probability of Superiority

The probability of superiority can be calculated using the examples below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```

cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
z <- ATE / p$sd_therapist
# standardize using therapist SD

```

```
pnorm(z / sqrt(2))
```

```
[1] 0.7364554
```

```
# raw ES  
pnorm(ATE / sqrt(2), sd = p$sd_therapist)
```

```
[1] 0.7364554
```

```
# cohen's d  
ICC <- p$sd_therapist^2/(p$sd_therapist^2 + p$sd_error^2)  
pnorm(  
  cohensd / sqrt(2),  
  sd = p$sd_therapist / tot_sd  
)
```

```
[1] 0.7364554
```

```
pnorm(  
  cohensd / sqrt(2),  
  sd = sqrt(ICC)  
)
```

```
[1] 0.7364554
```

```
# using simulation  
n <- 1e5  
cc <- rnorm(n, 0, p$sd_therapist)  
tx <- rnorm(n, ATE, p$sd_therapist)  
mean(tx > cc)
```

```
[1] 0.73714
```

## Table

Table 4 show the overlap measures for a range of treatment effects and ICCs.  $\sigma_e^2$  is constant, so only the treatment and therapist effects are varied.

```

icc <- c(0.01, 0.05, 0.1, 0.2)
d <- c(0.2, 0.5, 0.8)
sigma_error2 <- 1
sigma_u2 <- get_therapist_sd_from_icc(icc, sigma_error2)^2
grid <- expand_grid(d, sigma_u2)
grid <- grid %>%
  mutate(
    sigma_error2,
    icc = sigma_u2 / (sigma_u2 + sigma_error2),
    ate = d * sqrt((sigma_u2 + sigma_error2))
  )
map2_dfr(
  grid$d,
  grid$icc,
  function(d, icc) {
    data.frame(
      d = d,
      icc = icc,
      overlap = 2 * pnorm(-abs(d) / 2, sd = sqrt(icc)),
      u3 = pnorm(d, sd = sqrt(icc)),
      prob_superiority = pnorm(d / sqrt(2), sd = sqrt(icc))
    )
  }
) %>%
mutate(
  overlap = round(overlap * 100, 0),
  u3 = round(u3 * 100, 0),
  prob_superiority = round(prob_superiority, 2),
) %>%
left_join(
  select(grid, ate, d, icc, sigma_error2, sigma_u2),
  by = c("d", "icc")
) %>%
arrange(d, icc) %>%
rename(
  "Overlap (%)" = overlap,
  "$U_3$ (%)" = u3,
  "Pr. superiority" = prob_superiority,
  "Cohen's *d*" = d,
  "ICC" = icc,

```

```

"ATE" = ate,
"$\\sigma_e^2$" = sigma_error2,
"$\\sigma_u^2$" = sigma_u2
) %>%
kable(digits = 2)

```

Table 4: Interpreting Therapist Effects Using Overlap Measures

Cohen's $d$	ICC	Overlap (%)	$U_3$ (%)	Pr. superiority	ATE	$\sigma_e^2$	$\sigma_u^2$
0.2	0.01	32	98	0.92	0.20	1	0.01
0.2	0.05	65	81	0.74	0.21	1	0.05
0.2	0.10	75	74	0.67	0.21	1	0.11
0.2	0.20	82	67	0.62	0.22	1	0.25
0.5	0.01	1	100	1.00	0.50	1	0.01
0.5	0.05	26	99	0.94	0.51	1	0.05
0.5	0.10	43	94	0.87	0.53	1	0.11
0.5	0.20	58	87	0.79	0.56	1	0.25
0.8	0.01	0	100	1.00	0.80	1	0.01
0.8	0.05	7	100	0.99	0.82	1	0.05
0.8	0.10	21	99	0.96	0.84	1	0.11
0.8	0.20	37	96	0.90	0.89	1	0.25

## Confidence Intervals

To draw inferences about the population values we will compute confidence intervals for the therapist effects measures. The simplest way to construct CIs for the ICCs, and overlap measures is to use parametric bootstrapping or a Bayesian model. Here I will provide an example using Bayesian methods, but a parametric bootstrap using `lme4::bootMer` would look very similar.

For this example, I will increase the sample size to 40 therapists per treatment group. Robust inference regarding therapists effects will require large sample sizes.

```

p$sd_therapist_confounding <- 0
tot_sd <- sqrt(
  p$sd_error^2 + p$sd_therapist^2
)
p$ATE <- 0.2 * tot_sd
p$n2 <- 40

```



```

file_path <- "tmp/ci_dataset.rds"
if (file.exists(file_path)) {
  d <- read_rds(file = file_path)
} else {
  d <- simulate_data(p)
  write_rds(
    d,
    file = file_path,
    compress = "gz"
  )
}

```

The Bayesian multilevel model is fit using brms.

```

fit_tmp <- brm(
  y_random ~ Z + (1 | therapist_random),
  data = d,
  file = "tmp/fit_brms"
)
#' Calculate various posterior therapist effects measures
#'
#' @param the fitted brms model
calc_overlap_bayes <- function(fit) {
  draws <- as_draws_df(fit)[, 1:4]
  draws %>%
    mutate(
      sd_u = sd_therapist_random__Intercept,
      z = b_Z / sd_u,
      overlap = 2 * pnorm(-abs(z) / 2),
      u3 = pnorm(z),
      PS = pnorm(z / sqrt(2)),
      sd_e = sigma,
      icc = sd_u^2 / (sd_u^2 + sd_e^2)
    ) %>%
    select(
      tx = b_Z,
      sdu = sd_u,
      overlap,
      u3,
      PS,
      sde = sd_e,

```

```

      icc
    ) %>%
    summarize_all(
      list(
        "mean" = mean,
        "median" = median,
        "lwr" = ~ quantile(.x, 0.025),
        "upr" = ~ quantile(.x, 0.975)
      )
    ) %>%
    pivot_longer(
      everything(),
      names_to = c(".param", ".value"),
      names_pattern = "(.*)_(.*)"
    )
}

```

We then calculate all effect sizes and construct confidence intervals using the 2.5th and 97.5th percentiles of the posterior distributions, as shown in Table 5.

```

calc_overlap_bayes(fit_tmp) %>%
  rename(parameter = .param) %>%
  mutate(
    parameter = case_when(
      parameter == "PS" ~ "Prob. superiority",
      parameter == "icc" ~ "ICC",
      parameter == "overlap" ~ "Overlap",
      parameter == "sde" ~ "Error SD,  $\sigma_e$ ",
      parameter == "sdu" ~ "Therapist SD,  $\sigma_u$ ",
      parameter == "tx" ~ "Treatment effect  $\delta_1$ ",
      parameter == "u3" ~ "Cohen's  $U_{33}$ ",
      TRUE ~ parameter
    ),
    ci = paste0("[", round(lwr, 2), ", ", round(upr, 2), "]")
  ) %>%
  select(
    "Parameter" = parameter,
    "Est. (mean)" = mean,
    "Est. (median)" = median,
    "95% CI" = ci
  ) %>%

```

```
kable(digits = 2)
```

Table 5: 95% Confidence Intervals

Parameter	Est. (mean)	Est. (median)	95% CI
Treatment effect $\delta_1$	0.06	0.06	[-0.19, 0.32]
Therapist SD, $\sigma_u$	0.19	0.19	[0.02, 0.38]
Overlap	0.71	0.79	[0, 0.99]
Cohen's $U3$	0.62	0.64	[0.08, 1]
Prob. superiority	0.60	0.60	[0.16, 1]
Error SD, $\sigma_e$	1.55	1.55	[1.48, 1.63]
ICC	0.02	0.02	[0, 0.06]

### Simulation - Check Coverage

We will run a small simulation to check the frequentist coverage of the computed confidence intervals, i.e., do they include the true value 95% of the time?

```
#' Perform one confidence interval simulation
#
#' @param i an integer representing the current simulation index
#' @p a powerlmm study parameter object
#' @fit_tmp a brmsfit object
bayes_sim <- function(i, p, fit_tmp) {
  fit_b <- update(fit_tmp, newdata = simulate_data(p))
  fit_b %>%
    calc_overlap_bayes() %>%
    mutate(sim = i)
}
# check for cached simulation
file_path <- "tmp/simulation_overlap_bayes.rds"
if (file.exists(file_path)) {
  res_b <- read_rds(file = file_path)
} else {
  MAX_CORES <- as.numeric(Sys.getenv("MAX_CORES"))
  N_SIM <- as.numeric(Sys.getenv("N_SIM_CI"))
  cl <- parallel::makeCluster(MAX_CORES)
  parallel::clusterEvalQ(
    cl,
    expr = {
```

```

      library(powerlmm)
      library(brms)
      library(dplyr)
      library(tidyr)
    }
  )
  clusterExport(
    cl,
    c(
      "solve_for_mean",
      "get_therapist_sd_from_icc",
      "calc_overlap_bayes"
    )
  )
  res_b <- parallel::parLapply(
    cl,
    seq_len(ifelse(is.na(N_SIM), 1000, N_SIM)),
    bayes_sim,
    p = p,
    fit_tmp = fit_tmp
  )
  stopCluster(cl)
  res_b <- res_b %>% do.call(rbind, .)
  write_rds(
    res_b,
    file = file_path,
    compress = "gz"
  )
}

```

Table 6 summarizes the results from the 1000 simulations. Posterior means are used as estimates, relative bias is slightly lower for posterior medians.

```

res_b %>%
  rename(parameter = .param) %>%
  mutate(
    theta = case_when(
      parameter == "tx" ~ p$ATE,
      parameter == "sdu" ~ p$sd_therapist,
      parameter == "sde" ~ p$sd_error,
      parameter == "icc" ~ p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2),
    )
  )

```

```

        parameter == "overlap" ~ 2 * pnorm(-abs(p$ATE / p$sd_therapist) / 2),
        parameter == "u3" ~ pnorm(p$ATE / p$sd_therapist),
        parameter == "PS" ~ pnorm((p$ATE / p$sd_therapist) / sqrt(2)),
    ),
    coverage = lwr < theta & upr > theta
) %>%
group_by(parameter) %>%
summarize(
    M_est = mean(mean),
    Med_est = mean(median),
    sd_mean = sd(mean),
    coverage = mean(coverage),
    theta = unique(theta),
    lwr = quantile(mean, 0.025),
    upr = quantile(mean, 0.975)
) %>%
mutate(
    parameter = case_when(
        parameter == "PS" ~ "Prob. superiority",
        parameter == "icc" ~ "ICC",
        parameter == "overlap" ~ "Overlap",
        parameter == "sde" ~ "Error SD,  $\sigma_e$ ",
        parameter == "sdu" ~ "Therapist SD,  $\sigma_u$ ",
        parameter == "tx" ~ "Treatment effect  $\delta_1$ ",
        parameter == "u3" ~ "Cohen's  $U_{3<sub>3</sub>}$ ",
        TRUE ~ parameter
    ),
    rel_bias = (M_est - theta) / theta,
    rel_bias_med = (Med_est - theta) / theta
) %>%
select(
    Parameter = parameter,
    "Estimate" = M_est,
    "SD(Est.)" = sd_mean,
    "Coverage" = coverage,
    "True" = theta,
    "2.5th" = lwr,
    "97.5th" = upr,
    "Rel. bias (Est.)" = rel_bias,
) %>%
kable(digits = 2)

```

Table 6: Simulation Results for the 95% CIs

Parameter	Estimate	SD(Est.)	Coverage	True	2.5th	97.5th	Rel. bias (Est.)
Prob. superiority	0.74	0.08	0.95	0.74	0.58	0.88	0.00
ICC	0.05	0.02	0.94	0.05	0.02	0.08	0.03
Overlap	0.64	0.13	0.95	0.65	0.38	0.86	-0.03
Error SD, $\sigma_e$	1.50	0.03	0.95	1.50	1.45	1.56	0.00
Therapist SD, $\sigma_u$	0.34	0.06	0.95	0.34	0.22	0.45	-0.01
Treatment effect $\delta_1$	0.31	0.11	0.95	0.31	0.10	0.52	0.01
Cohen's $U3$	0.80	0.09	0.95	0.81	0.61	0.95	-0.01

Table 7 shows the power for the treatment effect and the probability of superiority.

```

res_b %>%
  filter(.param == "tx") %>%
  group_by(.param) %>%
  summarize(power = 1 - mean(lwr <= 0 & upr >= 0)) %>%
  rbind(
    res_b %>%
      filter(.param == "PS") %>%
      group_by(.param) %>%
      summarise(power = 1 - mean(lwr <= 0.5 & upr >= 0.5))
  ) %>%
  mutate(
    .param = case_when(
      .param == "PS" ~ "Prob. superiority",
      .param == "tx" ~ "Treatment effect  $\delta_1$ "
    )
  ) %>%
  rename(
    "Parameter" = .param,
    "Power" = power
  ) %>%
  kable(digits = 2)

```

Table 7: Simulation-based Power

Parameter	Power
Treatment effect $\delta_1$	0.81
Prob. superiority	0.81

Parameter	Power
-----------	-------