

A Causal Inference Perspective on Therapist Effects—Online Supplement

Kristoffer Magnusson

2023-06-18

Table of contents

Introduction	1
Therapist-Outcome Confounding Simulation	2
A More Intuitive Interpretation of Therapist Effects	12
Confidence Intervals	24

Introduction

This online supplement for *A Causal Inference Perspective on Therapist Effects* (Magnusson, 2023, in preparation) contains the code for all calculations, simulations, and figures presented in the article.

- [Preprint](#)
- [GitHub repository](#)
- [OSF repository](#)
- [Interactive visualization](#)

An HTML version of this document can also be viewed at <https://rpsychologist.github.io/causal-therapist-effects-paper>

```
knitr::opts_chunk$set(  
  message = FALSE,  
  warning = FALSE,  
  cache = TRUE  
)
```

```

library(ggplot2)
library(readr)
library(dplyr)
library(tidyr)
library(purrr)
library(powerlmm)
library(knitr)
library(parallel)
library(lme4)
library(lmerTest)
library(svglight)
library(brms)
library(parallel)

```

Therapist-Outcome Confounding Simulation

The functions below were used to simulate a data set with therapist effects and a therapist-outcome confounder.

```

#' Get the mean difference between prognostic groups
#'
#' @param sd the SD added to the therapist level due to the confounder
#' @param n2 the number of patients per therapist
#'
#' @returns the mean difference between the two binary groups
#' @examples
#' n2 <- 50
#' M <- solve_for_mean(0.075, n2)
#' # we solve for M in this
#' sd(c(rep(0, n2), rep(M, n2)))
solve_for_mean <- function(sd, n2) {
  2 * sqrt(sd^2 * (n2 * 2 - 1) / (n2 * 2))
}

#' Get the therapist SD from the ICC
#'
#' @param icc the ICC
#' @param sd_error the error SD
#'
#' @returns the random therapist SD

```

```

get_therapist_sd_from_icc <- function(icc, sd_error) {
  (sqrt(icc) * sd_error) / (sqrt(1 - icc))
}

#' Simulate confounded therapist effects
#'
#' @param n1 the number of patients per therapist
#' @param n2 the number of therapists per treatment group
#' @param sd_therapist the random therapist SD
#' @param sd_therapist_confounding the SD added to sd_therapist due to confounding
#' @param sd_error the random error SD
#' @param ATE average treatment effect
#'
#' @returns a data.frame
simulate_therapist_effect_confounding <- function(
  n1,
  n2,
  sd_therapist,
  sd_therapist_confounding,
  sd_error,
  ATE,
  ...) {
  tot_n <- 2 * n1 * n2
  b_pre_prognosis <- solve_for_mean(
    sd = sd_therapist_confounding,
    n2 = n2
  )
  Z <- rep(c(0, 1), each = n1 * n2)
  therapist_cc <- 1:n2
  therapist_tx <- (n2 + 1):(n2 * 2)
  therapist_effect_cc <- rnorm(n2, 0, sd_therapist)
  therapist_effect_tx <- rnorm(n2, 0, sd_therapist)
  therapist_effects <- c(therapist_effect_cc, therapist_effect_tx)
  error <- rnorm(tot_n, 0, sd_error)
  d <- data.frame(
    Z,
    pre_prognosis = rbinom(tot_n, 1, 0.5),
    therapist = NA,
    therapist_random = NA
  )
  d$therapist <- ifelse(

```

```

Z == 0,
# control
ifelse(
  d$pre_prognosis == 0,
  sample(
    x = 1:(n2 / 2),
    size = n1 * n2 / 2,
    replace = TRUE
  ),
  sample(
    x = (n2 / 2 + 1):(n2),
    size = n1 * n2 / 2,
    replace = TRUE
  )
),
# treatment
ifelse(
  d$pre_prognosis == 0,
  sample(
    x = 1:(n2 / 2),
    size = n1 * n2 / 2,
    replace = TRUE
  ),
  sample(
    x = (n2 / 2 + 1):(n2),
    size = n1 * n2 / 2,
    replace = TRUE
  )
) + n2
)
d$therapist_random[d$Z == 0] <- sample(
  x = therapist_cc,
  size = n1 * n2,
  replace = TRUE
)
d$therapist_random[d$Z == 1] <- sample(
  x = therapist_tx,
  size = n1 * n2,
  replace = TRUE
)
# non-random allocation of therapists

```

```

d$therapist_effect <- therapist_effects[d$therapist]
d$y <- 10 +
  d$Z * ATE +
  d$therapist_effect +
  error +
  d$pre_prognosis * b_pre_prognosis
# random allocation of therapists
d$therapist_effect_random <- therapist_effects[d$therapist_random]
d$y_random <- 10 +
  d$Z * ATE +
  d$therapist_effect_random +
  error +
  d$pre_prognosis * b_pre_prognosis
# required by powerlmm
d$time <- 0

d
}

```

To run the simulation, we pass the simulation function `simulate_therapist_effect_confounding` and use a custom `powerlmm` model to set up all the parameters.

```

ds <- study_design(custom = TRUE)
ICC <- 0.05
sd_error <- 1.5
sd_therapist <- sqrt((ICC * sd_error^2) / (1 - ICC))
# check ICC
sd_therapist^2 / (sd_therapist^2 + sd_error^2)

```

```
[1] 0.05
```

```

p <- powerlmm:::study_parameters.plcp_design_custom(
  design = ds,
  n1 = 20,
  n2 = 10,
  sd_therapist = sd_therapist,
  sd_error = sd_error,
  ATE = 0,
  data_gen = simulate_therapist_effect_confounding
)

```

```

p$sd_therapist_confounding <- get_therapist_sd_from_icc(
  icc = 0.075,
  sd_error = p$sd_error
)
# sd within groups
tot_sd <- sqrt(
  p$sd_error^2 + p$sd_therapist^2 + p$sd_therapist_confounding^2
)
# Cohen's d = 0.5
p$ATE <- 0.5 * tot_sd
p$tot_n <- 2 * p$n1 * p$n2
p$df <- 2 * p$n2 - 2

```

True parameter values

```

p$thetas_FE <- list(
  "(Intercept)" = 10,
  "Z" = p$ATE
)
p$thetas_RE <- list(
  "therapist_(Intercept)" = p$sd_therapist^2,
  "therapist_random_(Intercept)" = p$sd_therapist^2,
  "error" = p$sd_error^2
)

```

We then specify the statistical models.

```

f0 <- sim_formula(
  "y_random ~ Z + (1 | therapist_random)",
  test = "Z"
)
f1 <- sim_formula(
  "y_random ~ Z + pre_prognosis + (1 | therapist_random)",
  test = "Z"
)
f2 <- sim_formula(
  "y ~ Z + (1 | therapist)",
  test = "Z"
)
f3 <- sim_formula(
  "y ~ Z + pre_prognosis + (1 | therapist)",

```

```

      test = "Z"
    )
    f4 <- sim_formula(
      "y ~ Z",
      test = "Z"
    )
    f5 <- sim_formula(
      "y ~ Z + pre_prognosis",
      test = "Z"
    )
    f <- sim_formula_compare(
      "rand" = f0,
      "rand_adj" = f1,
      "confounding" = f2,
      "adjusted" = f3,
      "ignored" = f4,
      "ignored_adjusted" = f5
    )

```

Finally, we run the simulation and save the results.

```

# Load cache if it exists
file_path <- "tmp/simulation.rds"
if (file.exists(file_path)) {
  res <- read_rds(file = file_path)
} else {
  MAX_CORES <- as.numeric(Sys.getenv("MAX_CORES"))
  N_SIM <- as.numeric(Sys.getenv("N_SIM"))
  if (is.na(MAX_CORES)) MAX_CORES <- parallel::detectCores(logical = FALSE) - 1
  cl <- makeCluster(MAX_CORES)
  clusterExport(
    cl,
    c(
      "solve_for_mean",
      "get_therapist_sd_from_icc"
    )
  )
  res <- simulate(
    p,
    nsim = ifelse(is.na(N_SIM), 10000, N_SIM),
    cores = MAX_CORES,

```

```

        formula = f,
        cl = cl,
        satterthwaite = TRUE
    )
stopCluster(cl)
write_rds(
  res,
  file = file_path,
  compress = "gz"
)
}

```

Number of simulations: 10^4

Simulations results

We first summarize the treatment effects, and the result is shown in [Table 1](#).

```

summary(
  res,
  verbose = FALSE,
  para = "Z"
)$summary$summary$FE %>%
  mutate(
    SD_rel_bias = (M_se - SD_est) / SD_est
  ) %>%
  relocate(model, .before = 1) %>%
  select(
    -parameter,
    -Power,
    -Power_bw,
    "Model" = model,
    "Estimate" = M_est,
    "Rel. bias (SE)" = SD_rel_bias,
    "SD(Est.)" = SD_est,
    "Power" = Power_satt,
    "True" = theta
  ) %>%
  kable(digits = 2)

```


Table 1: Simulation results (treatment effects)

	Model	Estimate	True	M_se	SD(Est.)	Power	Rel. bias (SE)
2	rand	0.8	0.8	0.22	0.22	0.93	0.00
21	rand_adj	0.8	0.8	0.21	0.21	0.94	0.00
22	confounding	0.8	0.8	0.29	0.21	0.80	0.35
23	adjusted	0.8	0.8	0.21	0.21	0.94	0.00
24	ignored	0.8	0.8	0.16	0.22	0.99	-0.28
25	ignored_adjusted	0.8	0.8	0.15	0.22	0.99	-0.29

A summary of the estimated ICCs is shown in Table 2.

```
lapply(
  seq_along(res$res),
  function(i) {
    res$res[[i]]$RE %>%
      group_by(sim) %>%
      summarize(ICC = vcov[1] / sum(vcov)) %>%
      ungroup() %>%
      summarize(
        est_mean = mean(ICC),
        est_sd = sd(ICC),
        est_lwr = quantile(ICC, 0.025),
        est_upr = quantile(ICC, 0.975)
      ) %>%
      ungroup() %>%
      mutate(
        model = names(res$res)[[i]],
        parameter = "ICC",
        .before = 1
      )
  }
) %>%
  bind_rows() %>%
  filter(grepl("ignored", model) == FALSE) %>%
  mutate(
    theta = p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2),
    rel_bias = (est_mean - theta) / theta
  ) %>%
  select(
    "Model" = model,
```

```

    "Parameter" = parameter,
    "Estimate" = est_mean,
    "Rel. bias" = rel_bias,
    "SD(Est.)" = est_sd,
    "Est. (2.5%)" = est_lwr,
    "Est. (97.5%)" = est_upr,
    "True" = theta
  ) %>%
  kable(digits = 2)

```

Table 2: Simulation results (ICCs)

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand	ICC	0.05	-0.07	0.03	0.00	0.11	0.05
rand_adj	ICC	0.05	0.00	0.03	0.00	0.12	0.05
confounding	ICC	0.12	1.39	0.04	0.04	0.21	0.05
adjusted	ICC	0.05	-0.01	0.03	0.00	0.12	0.05

Lastly, we summarize the variance components.

```

lapply(
  seq_along(res$res),
  function(i) {
    res$res[[i]]$RE %>%
      group_by(parameter) %>%
      summarize(
        est_mean = mean(vcov),
        est_sd = sd(vcov),
        est_lwr = quantile(vcov, 0.025),
        est_upr = quantile(vcov, 0.975)
      ) %>%
      ungroup() %>%
      mutate(
        model = names(res$res)[[i]],
        .before = 1
      )
  }
) %>%
  bind_rows() %>%

```

```

left_join(
  data.frame(
    parameter = names(p$thetas_RE),
    theta = unlist(p$thetas_RE)
  )
) %>%
mutate(
  parameter = replace(
    parameter,
    grep("therapist", parameter),
    "therapist"
  ),
  theta = case_when(
    parameter == "therapist" ~ p$sd_therapist^2,
    parameter == "error" & model == "rand" ~ p$sd_error^2 +
      p$sd_therapist_confounding^2,
    model == "ignored" ~ p$sd_error^2 + p$sd_therapist^2 +
      p$sd_therapist_confounding^2,
    model == "ignored_adjusted" ~ p$sd_error^2 + p$sd_therapist^2,
    parameter == "error" ~ p$sd_error^2
  ),
  rel_bias = (est_mean - theta) / theta
) %>%
select(
  "Model" = model,
  "Parameter" = parameter,
  "Estimate" = est_mean,
  "Rel. bias" = rel_bias,
  "SD(Est.)" = est_sd,
  "Est. (2.5%)" = est_lwr,
  "Est. (97.5%)" = est_upr,
  "True" = theta
) %>%
kable(digits = 2)

```

Table 3: Simulation results (variance components)

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand	error	2.42	0.00	0.18	2.08	2.78	2.43
rand	therapist	0.12	0.01	0.08	0.00	0.30	0.12

Model	Parameter	Estimate	Rel. bias	SD(Est.)	Est. (2.5%)	Est. (97.5%)	True
rand_adj	error	2.25	0.00	0.16	1.93	2.58	2.25
rand_adj	therapist	0.12	0.01	0.08	0.00	0.30	0.12
confounding	error	2.25	0.00	0.17	1.93	2.58	2.25
confounding	therapist	0.31	1.61	0.13	0.10	0.59	0.12
adjusted	error	2.25	0.00	0.17	1.93	2.58	2.25
adjusted	therapist	0.12	0.00	0.08	0.00	0.30	0.12
ignored	error	2.53	-0.01	0.19	2.16	2.93	2.55
ignored_adj	error	2.35	-0.01	0.17	2.03	2.70	2.37

A More Intuitive Interpretation of Therapist Effects

This section presents R code for calculating the overlap measures presented in the manuscript.

Calculate the overlap between therapist distributions

Overlap can be calculated using the examples presented below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```
# Integration
int_f <- function(x, mu1, mu2, sd1, sd2) {
  f1 <- dnorm(x, mean = mu1, sd = sd1)
  f2 <- dnorm(x, mean = mu2, sd = sd2)
  pmin(f1, f2)
}

cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
# standardize using therapist SD
z <- ATE / p$sd_therapist
2 * pnorm(-abs(z) / 2)
```

```
[1] 0.6547208
```

```
# unstandardized
2 * pnorm(
  -abs(ATE) / 2,
```

```
sd = p$sd_therapist
)
```

```
[1] 0.6547208
```

```
# Integrate unstandardized
integrate(
  int_f,
  -Inf,
  Inf,
  mu1 = 0,
  mu2 = ATE,
  sd1 = p$sd_therapist,
  sd2 = p$sd_therapist
)
```

```
0.6547209 with absolute error < 1.6e-05
```

```
# cohen's d
2 * pnorm(
  -abs(cohensd) / 2,
  sd = p$sd_therapist / tot_sd
)
```

```
[1] 0.6547208
```

```
ICC <- p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2)
2 * pnorm(
  -abs(cohensd) / 2,
  sd = sqrt(ICC)
)
```

```
[1] 0.6547208
```

Plot Overlap

Figure 1 visualizes the overlapping therapist effect distributions.

```

SD <- sqrt(0.05)
mean1 <- 0.2
# create x axis
x_min <- 0 - 3 * SD
x_max <- mean1 + 3 * SD
x <- seq(x_min, x_max, length.out = 2e4)
df_control <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, 0, SD)),
  data.frame("x" = x_max, "y" = 0)
)
df_tx <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, mean1, SD)),
  data.frame("x" = x_max, "y" = 0)
)
poly_overlap <- data.frame(
  "x" = df_control$x,
  "y" = pmin(df_control$y, df_tx$y)
)
# colors
overlap_fill <- "#2980b9"
u3_fill <- "#3498db"
control_fill <- "#7f8c8d"
treatment_fill <- "#2c3e50"
p0 <- ggplot(
  df_tx,
  aes(
    x,
    y,
    fill = "treatment",
  )
) +
  # fill treatment group
  geom_polygon(
    linewidth = 1,
  ) +
  # fill control group
  geom_polygon(
    data = df_control,
    aes(

```

```

        fill = "control"
    ),
    linewidth = 1
) +
# overlap
geom_polygon(
    data = poly_overlap,
    color = NA,
    fill = overlap_fill,
) +
# line treatment
geom_polygon(
    linewidth = 1,
    color = "white",
    alpha = 0.5,
    fill = NA
) +
# line control
geom_polygon(
    data = df_control,
    linewidth = 1,
    color = "white",
    alpha = 0.5,
    fill = NA
) +
geom_vline(
    xintercept = 0,
    linetype = "dotted"
) +
geom_vline(
    xintercept = mean1,
    linetype = "dotted"
) +
annotate(
    geom = "text",
    label = "Control",
    x = 0,
    y = dnorm(0, 0, SD) * 1.1
) +
annotate(
    geom = "text",

```

```

    label = "Treatment",
    x = mean1,
    y = dnorm(mean1, mean1, SD) * 1.1
  ) +
  scale_color_manual(
    values = c(
      "control" = control_fill,
      "treatment" = treatment_fill
    )
  ) +
  scale_fill_manual(
    values = c(
      "control" = control_fill,
      "treatment" = treatment_fill
    )
  ) +
  labs(x = "Therapist effects", y = NULL) +
  theme_minimal() +
  theme(
    legend.position = "none",
    panel.grid.minor.y = element_blank(),
    panel.grid.major.y = element_blank(),
    axis.text.y = element_blank()
  )
p0
ggsave("figures/fig_overlap.svg", width = 8, height = 3)

```





Figure 1: The proportion of therapists from each treatment group with overlapping effects.

Cohen's U_3

Cohen's U_3 can be calculated using the examples below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```
# standardize using therapist SD
cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
# standardize using therapist SD
z <- ATE / p$sd_therapist
pnorm(z)
```

```
[1] 0.8144533
```

```
# unstandardized effect
pnorm(ATE, sd = p$sd_therapist)
```

```
[1] 0.8144533
```

```
# cohen's d parameterization
pnorm(
  cohensd,
  sd = p$sd_therapist / tot_sd
)
```

```
[1] 0.8144533
```

```
ICC <- p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2)
pnorm(
  cohensd,
  sd = sqrt(ICC)
)
```

```
[1] 0.8144533
```

Plot Cohen's U_3

Figure 2 visualizes Cohen's U_3 .

```
SD <- sqrt(0.05)
ES <- 0.2
mean1 <- ES
# create x axis
x_min <- 0 - 3 * SD
x_max <- mean1 + 3 * SD
x <- seq(x_min, x_max, length.out = 2e4)
df_control <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, 0, SD)),
  data.frame("x" = x_max, "y" = 0)
)
df_tx <- rbind(
  data.frame("x" = x_min, "y" = 0),
  data.frame("x" = x, "y" = dnorm(x, mean1, SD)),
  data.frame("x" = x_max, "y" = 0)
)
poly_u3 <- rbind(
  data.frame("x" = x_min, "y" = 0),
```

```

    poly_overlap[poly_overlap$x <= 0, ],
    data.frame("x" = 0, "y" = 0)
  )
  # colors
  overlap_fill <- "#2980b9"
  treatment_fill <- "#3498db"
  control_fill <- "#7f8c8d"
  u3_fill <- "#2c3e50"
  # plot
  p0 <- ggplot(
    df_control,
    aes(
      x,
      y,
      fill = "control"
    )
  ) +
    # fill control group
    geom_polygon(
      linewidth = 1,
    ) +
    # fill treatment group
    geom_polygon(
      data = df_tx,
      aes(
        fill = "treatment"
      ),
      linewidth = 1
    ) +
    # overlap
    geom_polygon(
      data = poly_u3,
      color = NA,
      fill = u3_fill
    ) +
    # line control
    geom_polygon(
      linewidth = 1,
      color = "white",
      alpha = 0.5,
      fill = NA
    )

```

```

) +
# line treatment
geom_polygon(
  data = df_tx,
  linewidth = 1,
  color = "white",
  alpha = 0.5,
  fill = NA
) +
geom_vline(
  xintercept = 0,
  linetype = "dotted"
) +
geom_vline(
  xintercept = mean1,
  linetype = "dotted"
) +
annotate(
  geom = "text",
  label = "Control",
  x = 0,
  y = dnorm(0, 0, SD) * 1.1
) +
annotate(
  geom = "text",
  label = "Treatment",
  x = mean1,
  y = dnorm(mean1, mean1, SD) * 1.1
) +
scale_color_manual(
  values = c(
    "control" = control_fill,
    "treatment" = treatment_fill
  )
) +
scale_fill_manual(
  values = c(
    "control" = control_fill,
    "treatment" = treatment_fill
  )
) +

```

```

labs(x = "Therapist effects", y = NULL) +
theme_minimal() +
theme(
  legend.position = "none",
  panel.grid.minor.y = element_blank(),
  panel.grid.major.y = element_blank(),
  axis.text.y = element_blank()
)
p0
ggsave("figures/fig_u3.svg", width = 8, height = 3)

```




Figure 2: A visualization of Cohen's U3, the proportion of the therapists in the treatment group with causal effects above the average therapist in the control group

Probability of Superiority

The probability of superiority can be calculated using the examples below, and several equivalent parameterizations are shown, using both standardized and raw effect sizes.

```

cohensd <- 0.2
tot_sd <- sqrt(p$sd_therapist^2 + p$sd_error^2)
ATE <- cohensd * tot_sd
z <- ATE / p$sd_therapist
# standardize using therapist SD

```

```
pnorm(z / sqrt(2))
```

```
[1] 0.7364554
```

```
# raw ES
pnorm(ATE / sqrt(2), sd = p$sd_therapist)
```

```
[1] 0.7364554
```

```
# cohen's d
ICC <- p$sd_therapist^2/(p$sd_therapist^2 + p$sd_error^2)
pnorm(
  cohensd / sqrt(2),
  sd = p$sd_therapist / tot_sd
)
```

```
[1] 0.7364554
```

```
pnorm(
  cohensd / sqrt(2),
  sd = sqrt(ICC)
)
```

```
[1] 0.7364554
```

```
# using simulation
n <- 1e5
cc <- rnorm(n, 0, p$sd_therapist)
tx <- rnorm(n, ATE, p$sd_therapist)
mean(tx > cc)
```

```
[1] 0.73645
```

Table

Table 4 show the overlap measures for a range of treatment effects and ICCs. σ_e^2 is constant, so only the treatment and therapist effects are varied.

```

icc <- c(0.01, 0.05, 0.1, 0.2)
d <- c(0.2, 0.5, 0.8)
sigma_error2 <- 1
sigma_u2 <- get_therapist_sd_from_icc(icc, sigma_error2)^2
grid <- expand_grid(d, sigma_u2)
grid <- grid %>%
  mutate(
    sigma_error2,
    icc = sigma_u2 / (sigma_u2 + sigma_error2),
    ate = d * sqrt((sigma_u2 + sigma_error2))
  )
map2_dfr(
  grid$d,
  grid$icc,
  function(d, icc) {
    data.frame(
      d = d,
      icc = icc,
      overlap = 2 * pnorm(-abs(d) / 2, sd = sqrt(icc)),
      u3 = pnorm(d, sd = sqrt(icc)),
      prob_superiority = pnorm(d / sqrt(2), sd = sqrt(icc))
    )
  }
) %>%
mutate(
  overlap = round(overlap * 100, 0),
  u3 = round(u3 * 100, 0),
  prob_superiority = round(prob_superiority, 2),
) %>%
left_join(
  select(grid, ate, d, icc, sigma_error2, sigma_u2),
  by = c("d", "icc")
) %>%
arrange(d, icc) %>%
rename(
  "Overlap (%)" = overlap,
  "$U_3$ (%)" = u3,
  "Pr. superiority" = prob_superiority,
  "Cohen's *d*" = d,
  "ICC" = icc,

```

```

"ATE" = ate,
"$\\sigma_e^2$" = sigma_error2,
"$\\sigma_u^2$" = sigma_u2
) %>%
kable(digits = 2)

```

Table 4: Interpreting Therapist Effects Using Overlap Measures

Cohen's d	ICC	Overlap (%)	U_3 (%)	Pr. superiority	ATE	σ_e^2	σ_u^2
0.2	0.01	32	98	0.92	0.20	1	0.01
0.2	0.05	65	81	0.74	0.21	1	0.05
0.2	0.10	75	74	0.67	0.21	1	0.11
0.2	0.20	82	67	0.62	0.22	1	0.25
0.5	0.01	1	100	1.00	0.50	1	0.01
0.5	0.05	26	99	0.94	0.51	1	0.05
0.5	0.10	43	94	0.87	0.53	1	0.11
0.5	0.20	58	87	0.79	0.56	1	0.25
0.8	0.01	0	100	1.00	0.80	1	0.01
0.8	0.05	7	100	0.99	0.82	1	0.05
0.8	0.10	21	99	0.96	0.84	1	0.11
0.8	0.20	37	96	0.90	0.89	1	0.25

Confidence Intervals

We will compute confidence intervals for the therapist effects measures to draw inferences about the population values. The simplest way to construct CIs for the ICCs, and overlap measures is to use parametric bootstrapping or a Bayesian model. I will provide an example using Bayesian methods, but a parametric bootstrap using `lme4::bootMer` would look very similar.

For this example, I will increase the sample size to 40 therapists per treatment group. Robust inference regarding therapist effects will require large sample sizes. Moreover, this is the required sample size to reach 80% power for $d = 0.2$.

```

p$sd_therapist_confounding <- 0
tot_sd <- sqrt(
  p$sd_error^2 + p$sd_therapist^2
)
p$ATE <- 0.2 * tot_sd

```



```

p$n2 <- 40
file_path <- "tmp/ci_dataset.rds"
if (file.exists(file_path)) {
  d <- read_rds(file = file_path)
} else {
  d <- simulate_data(p)
  write_rds(
    d,
    file = file_path,
    compress = "gz"
  )
}

```

The Bayesian multilevel model is fit using `brms`.

```

fit_tmp <- brm(
  y_random ~ Z + (1 | therapist_random),
  data = d,
  file = "tmp/fit_brms",
  refresh = 0
)

```

```
fit_tmp
```

```

Family: gaussian
Links: mu = identity; sigma = identity
Formula: y_random ~ Z + (1 | therapist_random)
Data: d (Number of observations: 1600)
Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
       total post-warmup draws = 4000

Group-Level Effects:
~therapist_random (Number of levels: 80)

```

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sd(Intercept)	0.42	0.06	0.32	0.54	1.00	1436	2334

```

Population-Level Effects:

```

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	10.05	0.09	9.88	10.22	1.00	2370	2611
Z	0.38	0.12	0.13	0.62	1.00	2700	2869

Family Specific Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	1.50	0.03	1.45	1.56	1.00	6942	3253

Draws were sampled using sampling(NUTS). For each parameter, Bulk_ESS and Tail_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

We then calculate all effect sizes and construct confidence intervals using the 2.5th and 97.5th percentiles of the posterior distributions, as shown in Table 5.

```
#' Calculate various posterior therapist effects measures
#'
#' @param the fitted brms model
calc_overlap_bayes <- function(fit) {
  draws <- as_draws_df(fit)[, 1:4]
  draws %>%
    mutate(
      sd_u = sd_therapist_random__Intercept,
      z = b_Z / sd_u,
      overlap = 2 * pnorm(-abs(z) / 2),
      u3 = pnorm(z),
      PS = pnorm(z / sqrt(2)),
      sd_e = sigma,
      icc = sd_u^2 / (sd_u^2 + sd_e^2)
    ) %>%
    select(
      tx = b_Z,
      sdu = sd_u,
      overlap,
      u3,
      PS,
      sde = sd_e,
      icc
    ) %>%
    summarize_all(
      list(
        "mean" = mean,
        "median" = median,
        "lwr" = ~ quantile(.x, 0.025),
        "upr" = ~ quantile(.x, 0.975)
      )
    )
}
```

```

    ) %>%
    pivot_longer(
      everything(),
      names_to = c(".param", ".value"),
      names_pattern = "(.*)_(.*)"
    )
  }

  calc_overlap_bayes(fit_tmp) %>%
    rename(parameter = .param) %>%
    mutate(
      parameter = case_when(
        parameter == "PS" ~ "Prob. superiority",
        parameter == "icc" ~ "ICC",
        parameter == "overlap" ~ "Overlap",
        parameter == "sde" ~ "Error SD,  $\sigma_e$ ",
        parameter == "sdu" ~ "Therapist SD,  $\sigma_u$ ",
        parameter == "tx" ~ "Treatment effect,  $\delta_1$ ",
        parameter == "u3" ~ "Cohen's  $U^3$ ",
        TRUE ~ parameter
      ),
      ci = paste0("[", round(lwr, 2), ", ", round(upr, 2), "]")
    ) %>%
    select(
      "Parameter" = parameter,
      "Est. (mean)" = mean,
      "Est. (median)" = median,
      "95% CI" = ci
    ) %>%
    kable(digits = 2)

```

Table 5: 95% Confidence Intervals

Parameter	Est. (mean)	Est. (median)	95% CI
Treatment effect, δ_1	0.38	0.39	[0.13, 0.62]
Therapist SD, σ_u	0.42	0.42	[0.32, 0.54]
Overlap	0.65	0.65	[0.44, 0.88]
Cohen's U^3	0.81	0.82	[0.62, 0.94]
Prob. superiority	0.74	0.74	[0.59, 0.86]
Error SD, σ_e	1.50	1.50	[1.45, 1.56]
ICC	0.07	0.07	[0.04, 0.12]

Simulation - Check Coverage

We will run a small simulation to check the frequentist coverage of the computed confidence intervals, i.e., do they include the true value 95% of the time?

```
#' Perform one confidence interval simulation
#'
#' @param i an integer representing the current simulation index
#' @p a powerlmm study parameter object
#' @fit_tmp a brmsfit object
bayes_sim <- function(i, p, fit_tmp) {
  fit_b <- update(fit_tmp, newdata = simulate_data(p))
  fit_b %>%
    calc_overlap_bayes() %>%
    mutate(sim = i)
}
# check for cached simulation
file_path <- "tmp/simulation_overlap_bayes.rds"
if (file.exists(file_path)) {
  res_b <- read_rds(file = file_path)
} else {
  MAX_CORES <- as.numeric(Sys.getenv("MAX_CORES"))
  N_SIM <- as.numeric(Sys.getenv("N_SIM_CI"))
  cl <- parallel::makeCluster(MAX_CORES)
  parallel::clusterEvalQ(
    cl,
    expr = {
      library(powerlmm)
      library(brms)
      library(dplyr)
      library(tidyr)
    }
  )
  clusterExport(
    cl,
    c(
      "solve_for_mean",
      "get_therapist_sd_from_icc",
      "calc_overlap_bayes"
    )
  )
  res_b <- parallel::parLapply(
```

```

      cl,
      seq_len(ifelse(is.na(N_SIM), 1000, N_SIM)),
      bayes_sim,
      p = p,
      fit_tmp = fit_tmp
    )
  stopCluster(cl)
  res_b <- res_b %>% do.call(rbind, .)
  write_rds(
    res_b,
    file = file_path,
    compress = "gz"
  )
}

```

Let's summarize the results from the 1000 simulations.

```

res_b %>%
  rename(parameter = .param) %>%
  mutate(
    theta = case_when(
      parameter == "tx" ~ p$ATE,
      parameter == "sdu" ~ p$sd_therapist,
      parameter == "sde" ~ p$sd_error,
      parameter == "icc" ~ p$sd_therapist^2 / (p$sd_therapist^2 + p$sd_error^2),
      parameter == "overlap" ~ 2 * pnorm(-abs(p$ATE / p$sd_therapist) / 2),
      parameter == "u3" ~ pnorm(p$ATE / p$sd_therapist),
      parameter == "PS" ~ pnorm((p$ATE / p$sd_therapist) / sqrt(2)),
    ),
    coverage = lwr < theta & upr > theta
  ) %>%
  group_by(parameter) %>%
  summarize(
    M_est = mean(mean),
    Med_est = mean(median),
    sd_mean = sd(mean),
    coverage = mean(coverage),
    theta = unique(theta),
    lwr = quantile(mean, 0.025),
    upr = quantile(mean, 0.975)
  ) %>%
  mutate(

```

```

parameter = case_when(
  parameter == "PS" ~ "Prob. superiority",
  parameter == "icc" ~ "ICC",
  parameter == "overlap" ~ "Overlap",
  parameter == "sde" ~ "Error SD,  $\sigma_e$ ",
  parameter == "sdu" ~ "Therapist SD,  $\sigma_u$ ",
  parameter == "tx" ~ "Treatment effect,  $\delta_1$ ",
  parameter == "u3" ~ "Cohen's  $U^3$ ",
  TRUE ~ parameter
),
rel_bias = (M_est - theta) / theta,
rel_bias_med = (Med_est - theta) / theta
) %>%
select(
  Parameter = parameter,
  "Estimate" = M_est,
  "SD(Est.)" = sd_mean,
  "Coverage" = coverage,
  "True" = theta,
  "2.5th" = lwr,
  "97.5th" = upr,
  "Rel. bias (Est.)" = rel_bias,
) %>%
kable(digits = 2)

```

Table 6: Simulation results (therapist effects)

Parameter	Estimate	SD(Est.)	Coverage	True	2.5th	97.5th	Rel. bias (Est.)
Prob. superiority	0.74	0.08	0.95	0.74	0.58	0.88	0.00
ICC	0.05	0.02	0.94	0.05	0.02	0.08	0.03
Overlap	0.64	0.13	0.95	0.65	0.38	0.86	-0.03
Error SD, σ_e	1.50	0.03	0.95	1.50	1.45	1.56	0.00
Therapist SD, σ_u	0.34	0.06	0.95	0.34	0.22	0.45	-0.01
Treatment effect, δ_1	0.31	0.11	0.95	0.31	0.10	0.52	0.01
Cohen's U^3	0.80	0.09	0.95	0.81	0.61	0.95	-0.01

Table 6 shows that all effects are unbiased and that the 95% CIs have nominal coverage. Posterior means are used as estimates; relative bias is slightly lower for posterior medians. It should be noted that the estimates will likely be biased with smaller sample sizes.

Table 7 shows the power for the treatment effect and the probability of superiority.

```

res_b %>%
  filter(.param == "tx") %>%
  group_by(.param) %>%
  summarize(power = 1 - mean(lwr <= 0 & upr >= 0)) %>%
  rbind(
    res_b %>%
      filter(.param == "PS") %>%
      group_by(.param) %>%
      summarise(power = 1 - mean(lwr <= 0.5 & upr >= 0.5))
  ) %>%
  mutate(
    .param = case_when(
      .param == "PS" ~ "Prob. superiority",
      .param == "tx" ~ "Treatment effect,  $\delta_1$ "
    )
  ) %>%
  rename(
    "Parameter" = .param,
    "Power" = power
  ) %>%
  kable(digits = 2)

```

Table 7: Simulation-based Power

Parameter	Power
Treatment effect, δ_1	0.81
Prob. superiority	0.81